

Documentation Technique Complète - Infrastructure MLOps Ansible

Table des Matières

1. Structure du projet
 2. Fichiers de configuration
 3. Variables globales
 4. Inventaire des serveurs
 5. Playbooks principaux
 6. Rôles Ansible détaillés
 7. Templates et configurations
 8. Commandes d'exécution
-

Structure du projet

Architecture complète du dossier ansible

```
/home/wassim/pfe/mlops/ansible/
  ansible.cfg                    # Configuration principale Ansible
  inventories/
    hosts.ini                    # Inventaire des 7 serveurs
  group_vars/
    all.yml                      # Variables pour tous les groupes
    haproxy.yml                  # Variables spécifiques HAProxy
    k8s-cluster.yml              # Variables cluster Kubernetes
  playbooks/
    site.yml                     # Playbook principal orchestrateur
    install-metallb.yml          # Installation MetallB
    install-prometheus.yml       # Installation stack monitoring
  roles/
    common/
      tasks/
        main.yml                 # Tâches communes tous serveurs
    haproxy/
      handlers/
        main.yml                 # Handlers restart/reload HAProxy
      tasks/
        main.yml                 # Installation et configuration HAProxy
      templates/
        haproxy.cfg.j2           # Template configuration HAProxy
    kubernetes-master/
      tasks/
        main.yml                 # Configuration masters Kubernetes
    kubernetes-worker/
```

```

        tasks/
            main.yml                    # Configuration workers Kubernetes
metallb/
    tasks/
        main.yml                    # Installation MetallB

```

Fichiers de configuration

ansible.cfg

```

[defaults]
# Configuration Ansible principale
inventory = inventories/hosts.ini
remote_user = user
private_key_file = ~/.ssh/id_rsa
host_key_checking = False
gather_facts = True
timeout = 30
retry_files_enabled = False

# Optimisation performances
forks = 10
gathering = smart
fact_caching = memory
fact_caching_timeout = 86400

# Logging
log_path = /tmp/ansible.log

[ssh_connection]
ssh_args = -o ControlMaster=auto -o ControlPersist=60s -o StrictHostKeyChecking=no
pipelining = True
control_path = /tmp/ansible-ssh-%%h-%%p-%%r

```

Explications :

- `inventory` : Pointe vers le fichier d'inventaire des serveurs
 - `remote_user` : Utilisateur pour connexions SSH (user)
 - `host_key_checking = False` : Désactive vérification clés SSH
 - `forks = 10` : Exécution parallèle sur 10 hôtes maximum
 - `gathering = smart` : Collecte intelligente des faits
 - `pipelining = True` : Optimisation SSH
-

Variables globales

group_vars/all.yml

```
---
# Variables globales pour tous les groupes d'hôtes

# Configuration utilisateur système
ansible_user: user
ansible_ssh_common_args: '-o StrictHostKeyChecking=no'

# Configuration réseau cluster
pod_network_cidr: "10.244.0.0/16"           # Réseau pods Flannel
service_cidr: "10.96.0.0/12"               # Réseau services Kubernetes
cluster_dns: "10.96.0.10"                  # IP DNS cluster

# Versions logicielles
kubernetes_version: "1.31.4"
containerd_version: "1.7.27"
flannel_version: "latest"

# Configuration monitoring
monitoring_namespace: "monitoring"
prometheus_nodeport: 30900
grafana_nodeport: 30300
alertmanager_nodeport: 30903

# Configuration ingress
ingress_namespace: "ingress-nginx"
ingress_http_nodeport: 32624
ingress_https_nodeport: 31316

# Paramètres système
disable_swap: true
enable_ipv4_forward: true
bridge_nf_call_iptables: true
```

Explications :

- pod_network_cidr : Plage IP pour les pods (compatible Flannel)
- service_cidr : Plage IP pour les services Kubernetes
- kubernetes_version : Version stable LTS utilisée
- *_nodeport : Ports d'accès externe aux services

group_vars/haproxy.yml

```
---
# Variables spécifiques au load balancer HAProxy
```

```
# Configuration HAProxy
haproxy_stats_enabled: true
haproxy_stats_port: 8404
haproxy_stats_uri: "/stats"
haproxy_stats_user: "admin"
haproxy_stats_password: "admin123"
```

```
# Backends Kubernetes API
k8s_api_backend_port: 6443
k8s_api_backends:
  - name: "master-1"
    address: "10.110.190.22"
    port: 6443
    check: "check"
  - name: "master-2"
    address: "10.110.190.23"
    port: 6443
    check: "check"
  - name: "master-3"
    address: "10.110.190.24"
    port: 6443
    check: "check"
```

```
# Backends Ingress HTTP
ingress_http_backends:
  - name: "worker-1"
    address: "10.110.190.25"
    port: 32624
    check: "check"
  - name: "worker-2"
    address: "10.110.190.26"
    port: 32624
    check: "check"
  - name: "worker-3"
    address: "10.110.190.27"
    port: 32624
    check: "check"
```

```
# Backends Ingress HTTPS
ingress_https_backends:
  - name: "worker-1"
    address: "10.110.190.25"
    port: 31316
    check: "check"
  - name: "worker-2"
```

```

        address: "10.110.190.26"
        port: 31316
        check: "check"
-   name: "worker-3"
    address: "10.110.190.27"
    port: 31316
    check: "check"

# Timeouts HAProxy
haproxy_connect_timeout: "4000ms"
haproxy_client_timeout: "50000ms"
haproxy_server_timeout: "50000ms"

```

Explications :

- haproxy_stats_* : Configuration interface de statistiques
- k8s_api_backends : Liste des masters pour équilibrage API
- ingress_*_backends : Liste des workers pour trafic Ingress
- Timeouts optimisés pour environnement production

group_vars/k8s-cluster.yml

```

---
# Variables cluster Kubernetes

# Identification cluster
k8s_cluster_name: "mlops-cluster"
k8s_version: "1.31.4"
k8s_api_server_port: 6443
k8s_node_port_range: "30000-32767"

# Point d'entrée control plane (HAProxy)
k8s_control_plane_endpoint: "10.110.190.21:6443"

# Runtime conteneurs
container_runtime: containerd
container_runtime_version: "1.7.27"

# Plugin réseau CNI
cni_plugin: flannel
flannel_version: "latest"

# Contrôleur Ingress
ingress_controller: nginx
ingress_http_nodeport: 32624
ingress_https_nodeport: 31316

```

```

# Packages Kubernetes avec versions pinned
kubernetes_packages:
  - kubelet=1.31.4-1.1
  - kubeadm=1.31.4-1.1
  - kubectl=1.31.4-1.1

# Configuration cluster kubeadm
cluster_configuration:
  apiVersion: kubeadm.k8s.io/v1beta3
  kind: ClusterConfiguration
  kubernetesVersion: "1.31.4"
  controlPlaneEndpoint: "10.110.190.21:6443"
  networking:
    podSubnet: "10.244.0.0/16"
    serviceSubnet: "10.96.0.0/12"

# Spécifications nœuds
master_node_count: 3
worker_node_count: 3

# Configuration services
kube_dns_ip: "10.96.0.10"
cluster_dns_domain: "cluster.local"

# Configuration système
os_distribution: "Ubuntu"
os_version: "24.04.2"
kernel_version: "6.8.0-53-generic"
architecture: "x86_64"

# État cluster actuel (documenté)
cluster_status: "running"
cluster_nodes_total: 6
master_nodes_count: 3
worker_nodes_count: 3
monitoring_coverage: "100%"
monitoring_targets_total: 75
monitoring_targets_up: 75

# Configuration containerd
containerd_config_file: "/etc/containerd/config.toml"
containerd_systemd_cgroup: true
containerd_runtime_type: "io.containerd.runc.v2"

# Configuration Flannel CNI
flannel_namespace: "kube-flannel"

```

```

flannel_configmap: "kube-flannel-cfg"
flannel_backend_type: "vxlan"
flannel_enable_nftables: false
flannel_mtu: 1450

# Configuration Ingress NGINX
ingress_namespace: "ingress-nginx"
ingress_controller_name: "ingress-nginx-controller"
ingress_service_type: "NodePort"
ingress_class_name: "nginx"

# Gestion certificats
cert_manager_namespace: "cert-manager"
enable_cert_manager: false

# Monitoring et logging
enable_metrics_server: true
enable_kubernetes_dashboard: false

# Politiques réseau
enable_network_policies: false
default_deny_all: false

# Configuration stockage
default_storage_class: "local-path"
enable_persistent_volumes: true

# Configuration sécurité
enable_pod_security_policy: false
enable_rbac: true

# Configuration API Server
api_server_port: 6443
api_server_secure_port: 6443
api_server_insecure_port: 0

# Scheduler et Controller Manager
scheduler_bind_address: "0.0.0.0"
controller_manager_bind_address: "0.0.0.0"

# Configuration etcd
etcd_data_dir: "/var/lib/etcd"
etcd_listen_client_urls: "https://127.0.0.1:2379"
etcd_listen_peer_urls: "https://127.0.0.1:2380"

```

Explications :

- `k8s_control_plane_endpoint` : Point d'entrée via HAProxy
 - `kubernetes_packages` : Versions pinned pour stabilité
 - `cluster_configuration` : Configuration kubeadm complète
 - Paramètres réseau optimisés pour production
-

Inventaire des serveurs

inventories/hosts.ini

```
# Inventaire infrastructure MLOps
# Réseau: 10.110.190.0/24

[haproxy]
haproxy-1 ansible_host=10.110.190.21 ansible_user=user hostname=haproxy

[k8s-masters]
master-1 ansible_host=10.110.190.22 ansible_user=user hostname=k8s-master-1
master-2 ansible_host=10.110.190.23 ansible_user=user hostname=k8s-master-2
master-3 ansible_host=10.110.190.24 ansible_user=user hostname=k8s-master-3

[k8s-workers]
worker-1 ansible_host=10.110.190.25 ansible_user=user hostname=k8s-worker-1
worker-2 ansible_host=10.110.190.26 ansible_user=user hostname=k8s-worker-2
worker-3 ansible_host=10.110.190.27 ansible_user=user hostname=k8s-worker-3

[k8s-cluster:children]
k8s-masters
k8s-workers

[all:vars]
ansible_ssh_common_args='-o StrictHostKeyChecking=no'
```

Structure :

- **1 HAProxy** : Load balancer externe (10.110.190.21)
 - **3 Masters** : Control plane haute disponibilité (10.110.190.22-24)
 - **3 Workers** : Nœuds de traitement (10.110.190.25-27)
 - **Groupe k8s-cluster** : Union masters + workers
-

Playbooks principaux

playbooks/site.yml

```
---
# Playbook principal orchestrateur infrastructure MLOps
```



```

- name: Configure HAProxy Load Balancer
  hosts: haproxy
  become: true
  roles:
    - haproxy
  tags: haproxy

- name: Configure Kubernetes Master Nodes
  hosts: k8s-masters
  become: true
  roles:
    - common
    - kubernetes-master
  tags: k8s-master

- name: Configure Kubernetes Worker Nodes
  hosts: k8s-workers
  become: true
  roles:
    - common
    - kubernetes-worker
  tags: k8s-worker

- name: Install Prometheus Monitoring Stack
  import_playbook: install-prometheus.yml
  tags: prometheus

- name: Install MetalLB Load Balancer
  import_playbook: install-metallb.yml
  tags: metallb

```

Flux d'exécution :

1. **HAProxy** : Configuration load balancer
2. **Masters** : Installation control plane Kubernetes
3. **Workers** : Ajout nœuds worker au cluster
4. **Prometheus** : Déploiement stack monitoring
5. **MetalLB** : Installation load balancer interne

playbooks/install-metallb.yml

```

---
# Installation MetalLB pour services LoadBalancer

- name: Install MetalLB Load Balancer
  hosts: k8s-masters[0]

```

```

become: true
roles:
  - metallb

post_tasks:
  - name: Verify MetalLB
    become_user: "{{ ansible_user }}"
    command: kubectl get pods -n metallb-system
    register: metallb_pods

  - name: Show MetalLB pods
    debug:
      var: metallb_pods.stdout_lines
    when: metallb_pods.stdout_lines is defined

  - name: Verify config
    become_user: "{{ ansible_user }}"
    command: kubectl get ipaddresspool -n metallb-system
    register: metallb_pool

  - name: Show IP pool
    debug:
      var: metallb_pool.stdout_lines
    when: metallb_pool.stdout_lines is defined

```

Fonctionnalité :

- Exécution sur premier master uniquement
- Installation MetalLB via manifests officiels
- Vérification post-installation
- Attente configuration IP pool par Zied

playbooks/install-prometheus.yml

```

---
# Installation stack monitoring Prometheus/Grafana

- name: Install kube-prometheus Stack
  hosts: k8s-masters[0]
  become: true
  tasks:
    - name: Create monitoring namespace
      become_user: "{{ ansible_user }}"
      command: kubectl create namespace {{ monitoring_namespace }}
      ignore_errors: true

    - name: Download kube-prometheus manifests

```

```

become_user: "{{ ansible_user }}"
get_url:
  url: "https://github.com/prometheus-operator/kube-prometheus/archive/refs/heads/main.zip"
  dest: "/tmp/kube-prometheus.zip"
  mode: '0644'

- name: Extract kube-prometheus
  become_user: "{{ ansible_user }}"
  unarchive:
    src: "/tmp/kube-prometheus.zip"
    dest: "/tmp/"
    remote_src: true

- name: Apply CRDs
  become_user: "{{ ansible_user }}"
  command: kubectl apply --server-side -f /tmp/kube-prometheus-main/manifests/setup/
  ignore_errors: true

- name: Wait for CRDs
  become_user: "{{ ansible_user }}"
  pause:
    seconds: 30

- name: Apply monitoring stack
  become_user: "{{ ansible_user }}"
  command: kubectl apply -f /tmp/kube-prometheus-main/manifests/
  ignore_errors: true

- name: Expose Prometheus NodePort
  become_user: "{{ ansible_user }}"
  command: |
    kubectl patch svc prometheus-k8s -n {{ monitoring_namespace }} -p '{
      "spec": {
        "type": "NodePort",
        "ports": [
          {
            "name": "web",
            "port": 9090,
            "targetPort": 9090,
            "nodePort": {{ prometheus_nodeport }}
          }
        ]
      }
    }'
  ignore_errors: true

```

```

- name: Expose Grafana NodePort
  become_user: "{{ ansible_user }}"
  command: |
    kubectl patch svc grafana -n {{ monitoring_namespace }} -p '{
      "spec": {
        "type": "NodePort",
        "ports": [
          {
            "name": "http",
            "port": 3000,
            "targetPort": 3000,
            "nodePort": {{ grafana_nodeport }}
          }
        ]
      }
    }'
  ignore_errors: true

- name: Verify monitoring deployment
  become_user: "{{ ansible_user }}"
  command: kubectl get pods -n {{ monitoring_namespace }}
  register: monitoring_pods

- name: Show monitoring pods
  debug:
    var: monitoring_pods.stdout_lines
  when: monitoring_pods.stdout_lines is defined

```

Composants installés :

- **Prometheus Operator** : Gestion déclarative monitoring
- **Prometheus** : Base de données métriques
- **Grafana** : Interface visualisation
- **AlertManager** : Gestion alertes
- **Node Exporter** : Métriques système
- **kube-state-metrics** : Métriques Kubernetes

Rôles Ansible détaillés

roles/common/tasks/main.yml

```

---
# Tâches communes à tous les serveurs

- name: Update system packages
  apt:

```

```

    update_cache: true
    upgrade: dist
    autoremove: true
    autoclean: true

- name: Install essential packages
  apt:
    name:
      - curl
      - wget
      - vim
      - htop
      - net-tools
      - software-properties-common
      - apt-transport-https
      - ca-certificates
      - gnupg
      - lsb-release
    state: present

- name: Disable swap permanently
  lineinfile:
    path: /etc/fstab
    regexp: '^.*swap.*$'
    state: absent
  when: disable_swap | default(true)

- name: Disable swap immediately
  command: swapoff -a
  when: disable_swap | default(true)

- name: Enable IPv4 forwarding
  sysctl:
    name: net.ipv4.ip_forward
    value: '1'
    state: present
    reload: true
  when: enable_ipv4_forward | default(true)

- name: Enable bridge netfilter
  sysctl:
    name: "{{ item }}"
    value: '1'
    state: present
    reload: true
  loop:

```

```

- net.bridge.bridge-nf-call-iptables
- net.bridge.bridge-nf-call-ip6tables
when: bridge_nf_call_iptables | default(true)

- name: Load kernel modules
modprobe:
  name: "{{ item }}"
  state: present
loop:
  - overlay
  - br_netfilter

- name: Make kernel modules persistent
lineinfile:
  path: /etc/modules-load.d/k8s.conf
  line: "{{ item }}"
  create: true
loop:
  - overlay
  - br_netfilter

- name: Set timezone
timezone:
  name: Europe/Paris

- name: Configure NTP
apt:
  name: ntp
  state: present

- name: Start and enable NTP
systemd:
  name: ntp
  state: started
  enabled: true

```

Préparation système :

- Mise à jour packages système
- Installation outils essentiels
- Désactivation swap (requis Kubernetes)
- Configuration réseau (IP forwarding, bridge netfilter)
- Chargement modules kernel
- Synchronisation temps

roles/haproxy/tasks/main.yml

```
---
# Installation et configuration HAProxy

- name: Install HAProxy
  apt:
    name: haproxy
    state: present
    update_cache: true

- name: Backup original HAProxy config
  copy:
    src: /etc/haproxy/haproxy.cfg
    dest: /etc/haproxy/haproxy.cfg.orig
    remote_src: true
    backup: true
    ignore_errors: true

- name: Generate HAProxy configuration
  template:
    src: haproxy.cfg.j2
    dest: /etc/haproxy/haproxy.cfg
    backup: true
    owner: root
    group: root
    mode: '0644'
  notify:
    - restart haproxy

- name: Start and enable HAProxy
  systemd:
    name: haproxy
    state: started
    enabled: true

- name: Verify HAProxy is running
  systemd:
    name: haproxy
    state: started
  register: haproxy_status

- name: Show HAProxy status
  debug:
    msg: "HAProxy is {{ haproxy_status.status.ActiveState }}"
```

```

- name: Test HAProxy configuration
  command: haproxy -c -f /etc/haproxy/haproxy.cfg
  register: haproxy_config_test
  changed_when: false

- name: Show HAProxy config test result
  debug:
    var: haproxy_config_test.stdout_lines

```

Installation HAProxy :

- Installation package HAProxy
- Sauvegarde configuration originale
- Génération configuration depuis template
- Démarrage et activation service
- Tests de validation

roles/haproxy/handlers/main.yml

```

---
# Handlers pour redémarrage HAProxy

- name: restart haproxy
  systemd:
    name: haproxy
    state: restarted
  listen: restart haproxy

- name: reload haproxy
  systemd:
    name: haproxy
    state: reloaded
  listen: reload haproxy

```

roles/kubernetes-master/tasks/main.yml

```

---
# Configuration nœuds masters Kubernetes

- name: Install containerd.io
  apt:
    name: containerd.io
    state: present
    update_cache: true

- name: Create containerd configuration directory
  file:

```



```

    path: /etc/containerd
    state: directory
    mode: '0755'

- name: Generate containerd default configuration
  shell: containerd config default > /etc/containerd/config.toml
  args:
    creates: /etc/containerd/config.toml

- name: Configure containerd to use systemd cgroup
  lineinfile:
    path: /etc/containerd/config.toml
    regexp: '^s*SystemdCgroup\s*='
    line: '          SystemdCgroup = true'
    backup: true

- name: Start and enable containerd
  systemd:
    name: containerd
    state: started
    enabled: true

- name: Add Kubernetes apt key
  apt_key:
    url: https://pkgs.k8s.io/core:/stable:/v1.31/deb/Release.key
    state: present

- name: Add Kubernetes repository
  apt_repository:
    repo: "deb https://pkgs.k8s.io/core:/stable:/v1.31/deb//"
    state: present
    filename: kubernetes

- name: Install specific Kubernetes components v1.31.4
  apt:
    name:
      - kubelet=1.31.4-1.1
      - kubeadm=1.31.4-1.1
      - kubectl=1.31.4-1.1
    state: present
    update_cache: true

- name: Hold Kubernetes packages
  dpkg_selections:
    name: "{{ item }}"
    selection: hold

```

```

loop:
  - kubelet
  - kubeadm
  - kubectl

- name: Start and enable kubelet
  systemd:
    name: kubelet
    state: started
    enabled: true

- name: Initialize Kubernetes cluster (only on first master)
  command: >
    kubeadm init
    --apiserver-advertise-address={{ ansible_host }}
    --pod-network-cidr={{ pod_network_cidr }}
    --service-cidr={{ service_cidr }}
    --control-plane-endpoint={{ k8s_control_plane_endpoint }}
    --upload-certs
    --kubernetes-version=v{{ k8s_version }}
  when: inventory_hostname == groups['k8s-masters'][0]
  register: kubeadm_init
  ignore_errors: true

- name: Generate join token for cluster
  command: kubeadm token create --print-join-command
  when: inventory_hostname == groups['k8s-masters'][0]
  register: join_command
  ignore_errors: true

- name: Generate certificate key for control plane
  command: kubeadm init phase upload-certs --upload-certs
  when: inventory_hostname == groups['k8s-masters'][0]
  register: certificate_key_output
  ignore_errors: true

- name: Extract certificate key
  set_fact:
    certificate_key: "{{ certificate_key_output.stdout_lines[-1] }}"
  when:
    - inventory_hostname == groups['k8s-masters'][0]
    - certificate_key_output.stdout_lines is defined
    - certificate_key_output.stdout_lines | length > 0

- name: Extract join command components
  set_fact:

```

```

    kubeadm_token: "{{ join_command.stdout.split('--token ')[1].split(' ')[0] }}"
    ca_cert_hash: "{{ join_command.stdout.split('--discovery-token-ca-cert-hash ')[1].split(' ')[0] }}"
when:
    - inventory_hostname == groups['k8s-masters'][0]
    - join_command.stdout is defined
    - "'--token ' in join_command.stdout"
    - "'--discovery-token-ca-cert-hash ' in join_command.stdout"

- name: Save join information for future reference
  copy:
    content: |
      # Kubernetes Cluster Join Information
      # Generated on: {{ ansible_date_time.date }} {{ ansible_date_time.time }}

      # Control Plane Endpoint
      CONTROL_PLANE_ENDPOINT={{ k8s_control_plane_endpoint }}

      # Join Token (expires in 24 hours by default)
      KUBEADM_TOKEN={{ kubeadm_token | default('FAILED_TO_GENERATE') }}

      # CA Certificate Hash
      CA_CERT_HASH={{ ca_cert_hash | default('FAILED_TO_GENERATE') }}

      # Certificate Key (for control plane nodes)
      CERTIFICATE_KEY={{ certificate_key | default('FAILED_TO_GENERATE') }}

      # Worker Join Command:
      # kubeadm join {{ k8s_control_plane_endpoint }} --token {{ kubeadm_token | default('FAILED_TO_GENERATE') }}

      # Master Join Command:
      # kubeadm join {{ k8s_control_plane_endpoint }} --token {{ kubeadm_token | default('FAILED_TO_GENERATE') }}
    dest: "/tmp/kubeadm-join-info.txt"
    mode: '0600'
  when: inventory_hostname == groups['k8s-masters'][0]

- name: Join additional master nodes to cluster
  command: >
    kubeadm join {{ k8s_control_plane_endpoint }}
    --token {{ hostvars[groups['k8s-masters'][0]]['kubeadm_token'] }}
    --discovery-token-ca-cert-hash {{ hostvars[groups['k8s-masters'][0]]['ca_cert_hash'] }}
    --control-plane
    --certificate-key {{ hostvars[groups['k8s-masters'][0]]['certificate_key'] }}
    --apiserver-advertise-address={{ ansible_host }}
  when: inventory_hostname != groups['k8s-masters'][0]
  register: kubeadm_join_master
  ignore_errors: true

```

```

- name: Create .kube directory for user
  file:
    path: /home/{{ ansible_user }}/.kube
    state: directory
    owner: "{{ ansible_user }}"
    group: "{{ ansible_user }}"
    mode: '0755'

- name: Copy admin.conf to user's kube config
  copy:
    src: /etc/kubernetes/admin.conf
    dest: /home/{{ ansible_user }}/.kube/config
    remote_src: true
    owner: "{{ ansible_user }}"
    group: "{{ ansible_user }}"
    mode: '0644'
  when: inventory_hostname == groups['k8s-masters'][0]

- name: Install Flannel CNI (only on first master)
  become_user: "{{ ansible_user }}"
  command: kubectl apply -f https://github.com/flannel-io/flannel/releases/latest/download/
  when: inventory_hostname == groups['k8s-masters'][0]
  ignore_errors: true

- name: Install NGINX Ingress Controller (only on first master)
  become_user: "{{ ansible_user }}"
  command: kubectl apply -f https://raw.githubusercontent.com/kubernetes/ingress-nginx/cont
  when: inventory_hostname == groups['k8s-masters'][0]
  ignore_errors: true

- name: Wait for ingress controller to be ready
  become_user: "{{ ansible_user }}"
  command: kubectl wait --namespace ingress-nginx --for=condition=ready pod --selector=app.k
  when: inventory_hostname == groups['k8s-masters'][0]
  ignore_errors: true

```

Configuration masters :

- Installation containerd avec systemd cgroup
- Installation Kubernetes v1.31.4 (version pinned)
- Initialisation cluster avec endpoint HAProxy
- Génération tokens pour join workers/masters
- Installation Flannel CNI et NGINX Ingress

roles/kubernetes-worker/tasks/main.yml

```
---
# Configuration nœuds workers Kubernetes

- name: Install containerd.io
  apt:
    name: containerd.io
    state: present
    update_cache: true

- name: Create containerd configuration directory
  file:
    path: /etc/containerd
    state: directory
    mode: '0755'

- name: Generate containerd default configuration
  shell: containerd config default > /etc/containerd/config.toml
  args:
    creates: /etc/containerd/config.toml

- name: Configure containerd to use systemd cgroup
  lineinfile:
    path: /etc/containerd/config.toml
    regexp: '^s*SystemdCgroup\s*='
    line: '          SystemdCgroup = true'
    backup: true

- name: Start and enable containerd
  systemd:
    name: containerd
    state: started
    enabled: true

- name: Add Kubernetes apt key
  apt_key:
    url: https://pkgs.k8s.io/core:/stable:/v1.31/deb/Release.key
    state: present

- name: Add Kubernetes repository
  apt_repository:
    repo: "deb https://pkgs.k8s.io/core:/stable:/v1.31/deb/ ."
    state: present
    filename: kubernetes
```

```

- name: Install specific Kubernetes components v1.31.4
  apt:
    name:
      - kubelet=1.31.4-1.1
      - kubeadm=1.31.4-1.1
    state: present
    update_cache: true

- name: Hold Kubernetes packages
  dpkg_selections:
    name: "{{ item }}"
    selection: hold
  loop:
    - kubelet
    - kubeadm

- name: Start and enable kubelet
  systemd:
    name: kubelet
    state: started
    enabled: true

- name: Join worker nodes to cluster
  command: >
    kubeadm join {{ k8s_control_plane_endpoint }}
    --token {{ hostvars[groups['k8s-masters'][0]]['kubeadm_token'] }}
    --discovery-token-ca-cert-hash {{ hostvars[groups['k8s-masters'][0]]['ca_cert_hash'] }}
  register: kubeadm_join_worker
  ignore_errors: true

- name: Show worker join result
  debug:
    var: kubeadm_join_worker.stdout_lines
    when: kubeadm_join_worker.stdout_lines is defined

```

Configuration workers :

- Installation containerd identique aux masters
- Installation kubelet et kubeadm (pas kubectl)
- Join automatique au cluster via HAProxy endpoint

roles/metallb/tasks/main.yml

```

---
# Installation MetallLB pour services LoadBalancer

- name: Install MetallLB namespace

```

```

become_user: "{{ ansible_user }}"
command: kubectl apply -f https://raw.githubusercontent.com/metallb/metallb/v0.14.8/config
when: inventory_hostname == groups['k8s-masters'][0]
register: metallb_install

- name: Wait for MetalLB controller
  become_user: "{{ ansible_user }}"
  command: kubectl wait --namespace metallb-system --for=condition=ready pod --selector=app=metallb
  when: inventory_hostname == groups['k8s-masters'][0]

- name: Show MetalLB status
  debug:
    msg: "MetalLB installed"
  when: inventory_hostname == groups['k8s-masters'][0]

```

Installation MetalLB :

- Déploiement manifests officiels v0.14.8
- Attente disponibilité pods MetalLB
- Configuration IP pool en attente de Zied

Templates et configurations

roles/haproxy/templates/haproxy.cfg.j2

Configuration HAProxy pour cluster Kubernetes MLOps
Généré automatiquement par Ansible

```

global
    log                  127.0.0.1:514 local0
    chroot               /var/lib/haproxy
    stats               socket /run/haproxy/admin.sock mode 660 level admin
    stats               timeout 30s
    user                haproxy
    group               haproxy
    daemon

defaults
    mode                tcp
    log                 global
    option              tcplog
    option              dontlognull
    option              redispatch
    retries             3
    timeout connect     {{ haproxy_connect_timeout | default('4000ms') }}
    timeout client      {{ haproxy_client_timeout | default('50000ms') }}

```

```

        timeout server          {{ haproxy_server_timeout | default('50000ms') }}
        maxconn                  3000

# Interface de statistiques HAProxy
{% if haproxy_stats_enabled | default(true) %}
listen stats
    bind *:{{ haproxy_stats_port | default(8404) }}
    mode http
    stats enable
    stats uri {{ haproxy_stats_uri | default('/stats') }}
    stats realm HAProxy\ Statistics
    stats auth {{ haproxy_stats_user | default('admin') }}:{{ haproxy_stats_password | default('') }}
    stats refresh 30s
    stats show-node
    stats show-legends
{% endif %}

# Frontend pour API Kubernetes (port 6443)
frontend k8s-api-frontend
    bind *:{{ k8s_api_server_port | default(6443) }}
    mode tcp
    option tcplog
    default_backend k8s-api-backend

# Backend API Kubernetes - Load balancing vers masters
backend k8s-api-backend
    mode tcp
    option tcp-check
    balance roundrobin
{% for backend in k8s_api_backends %}
    server {{ backend.name }} {{ backend.address }}:{{ backend.port }} {{ backend.check }}
{% endfor %}

# Frontend pour trafic HTTP Ingress (port 80)
frontend ingress-http-frontend
    bind *:80
    mode tcp
    option tcplog
    default_backend ingress-http-backend

# Backend HTTP Ingress - Load balancing vers workers
backend ingress-http-backend
    mode tcp
    option tcp-check
    balance roundrobin
{% for backend in ingress_http_backends %}

```



```

        server {{ backend.name }} {{ backend.address }}:{{ backend.port }} {{ backend.check }}
    {% endfor %}

# Frontend pour trafic HTTPS Ingress (port 443)
frontend ingress-https-frontend
    bind *:443
    mode tcp
    option tcplog
    default_backend ingress-https-backend

# Backend HTTPS Ingress - Load balancing vers workers
backend ingress-https-backend
    mode tcp
    option tcp-check
    balance roundrobin
{% for backend in ingress_https_backends %}
    server {{ backend.name }} {{ backend.address }}:{{ backend.port }} {{ backend.check }}
{% endfor %}

```

Configuration HAProxy :

- **Frontend k8s-api** : Port 6443 vers masters
- **Frontend ingress-http** : Port 80 vers workers NodePort 32624
- **Frontend ingress-https** : Port 443 vers workers NodePort 31316
- **Health checks** : TCP avec fall/rise automatique
- **Stats interface** : Port 8404 avec auth admin/admin123

Commandes d'exécution

Déploiement complet

```

# Navigation vers le dossier
cd /home/wassim/pfe/mlops/ansible

# Vérification connectivité
ansible all -i inventories/hosts.ini -m ping

# Déploiement infrastructure complète
ansible-playbook -i inventories/hosts.ini playbooks/site.yml

# Déploiement avec logs verbeux
ansible-playbook -i inventories/hosts.ini playbooks/site.yml -v

```

Déploiement par composant

```
# HAProxy uniquement
ansible-playbook -i inventories/hosts.ini playbooks/site.yml --tags haproxy

# Masters Kubernetes uniquement
ansible-playbook -i inventories/hosts.ini playbooks/site.yml --tags k8s-master

# Workers Kubernetes uniquement
ansible-playbook -i inventories/hosts.ini playbooks/site.yml --tags k8s-worker

# Stack monitoring Prometheus uniquement
ansible-playbook -i inventories/hosts.ini playbooks/site.yml --tags prometheus

# MetalLB uniquement (comme demandé par Zied)
ansible-playbook -i inventories/hosts.ini playbooks/site.yml --tags metallb
```

Commandes de vérification

```
# État du cluster
ssh user@10.110.190.22
kubectl get nodes -o wide
kubectl get pods --all-namespaces

# Services et endpoints
kubectl get svc --all-namespaces
kubectl get endpoints --all-namespaces

# État MetalLB
kubectl get pods -n metallb-system
kubectl get ipaddresspool -n metallb-system

# État monitoring
kubectl get pods -n monitoring
kubectl get svc -n monitoring

# HAProxy stats
curl http://10.110.190.21:8404/stats
```

Debug et logs

```
# Logs Ansible
tail -f /tmp/ansible.log

# Test connectivité spécifique
ansible haproxy -i inventories/hosts.ini -m ping
ansible k8s-masters -i inventories/hosts.ini -m ping
```

```
ansible k8s-workers -i inventories/hosts.ini -m ping

# Dry-run pour test
ansible-playbook -i inventories/hosts.ini playbooks/site.yml --check

# Limitation à un groupe
ansible-playbook -i inventories/hosts.ini playbooks/site.yml --limit haproxy
```

Commandes post-installation

```
# Configuration kubeconfig local
ssh user@10.110.190.22
sudo cp /etc/kubernetes/admin.conf ~/.kube/config
sudo chown $(id -u):$(id -g) ~/.kube/config

# Test accès cluster via HAProxy
kubectl --server=https://10.110.190.21:6443 get nodes

# Vérification services externes
kubectl get svc --all-namespaces --field-selector spec.type=LoadBalancer
kubectl get svc --all-namespaces --field-selector spec.type=NodePort
```

Documentation technique complète couvrant tous les fichiers du dossier ansible avec explications détaillées de chaque composant, variable et configuration.

Généré le 2 juillet 2025 - Version 1.0